

## **IBM NAAN MUDHALVAN**

### **CLOUD APPLICATION DEVELOPMENT**

**PROJECT** : Data Warehousing with IBM Cloud Db2Warehouse

#### **Phase 4: Development Part 2**

In this part you will continue building your project.

Continue building the data warehouse by implementing ETL processes and enabling data exploration.

Implement ETL processes to extract, transform, and load data into the data warehouse.

Enable data architects to explore and analyze data within Db2 Warehouse using SQL queries and analysis techniques.

In Phase 4 of project, need to focus on implementing ETL (Extract, Transform, Load) processes to populate your data warehouse, which appears to be IBM Db2 Warehouse, and enable data architects to explore and analyze data using SQL queries and analysis techniques.

#### **ETL**

ETL, which stands for extract, transform and load, is a data integration process that combines data from multiple data sources into a single, consistent data store that is loaded into a [data warehouse](#) or other target system.

As the databases grew in popularity in the 1970s, ETL was introduced as a process for integrating and loading data for computation and analysis, eventually becoming the primary method to process data for data warehousing projects.

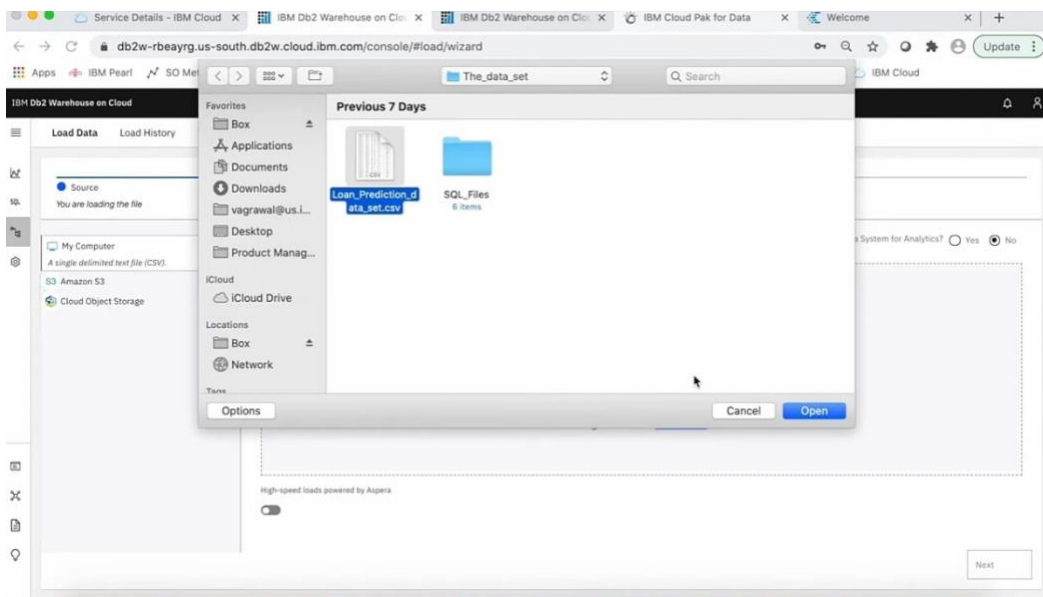
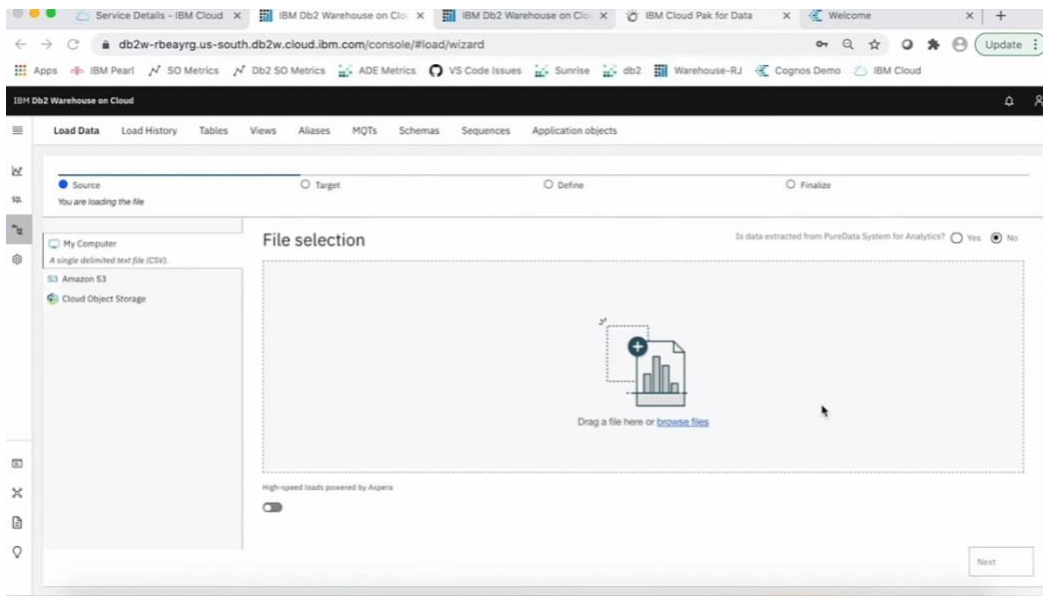
ETL provides the foundation for data analytics and machine learning workstreams. Through a series of business rules, ETL cleanses and organizes data in a way which addresses specific business intelligence needs, like monthly reporting, but it can also tackle more advanced analytics, which can improve back-end processes or end user experiences. ETL is often used by an organization to:

- Extract data from legacy systems
- Cleanse the data to improve data quality and establish consistency
- Load data into a target database

Here are some key steps you might consider:

## 1. Data Extraction:

Extract data from various source systems. This could involve databases, files, APIs, or other data repositories.

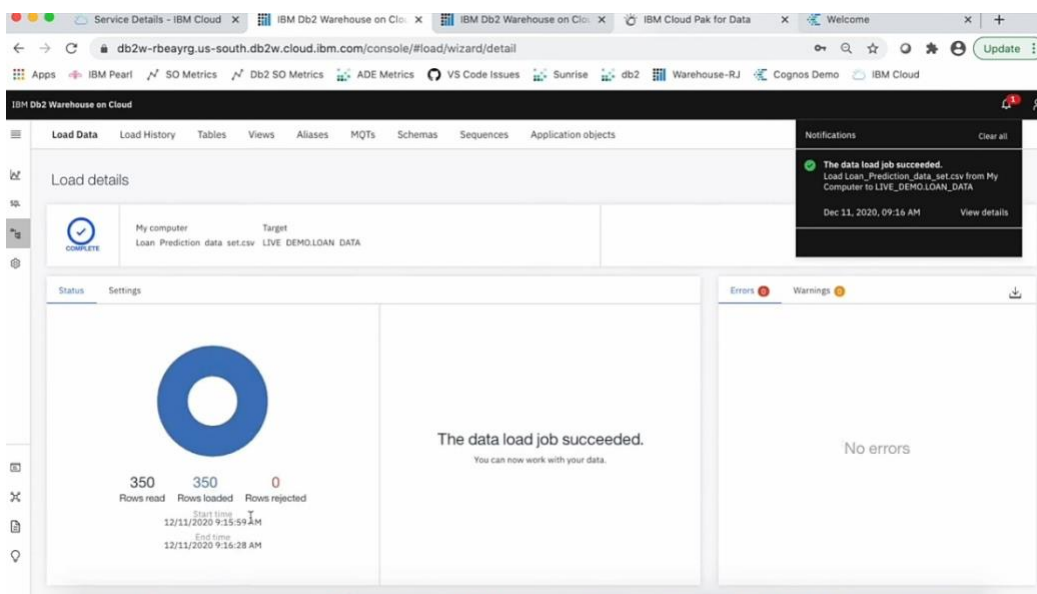
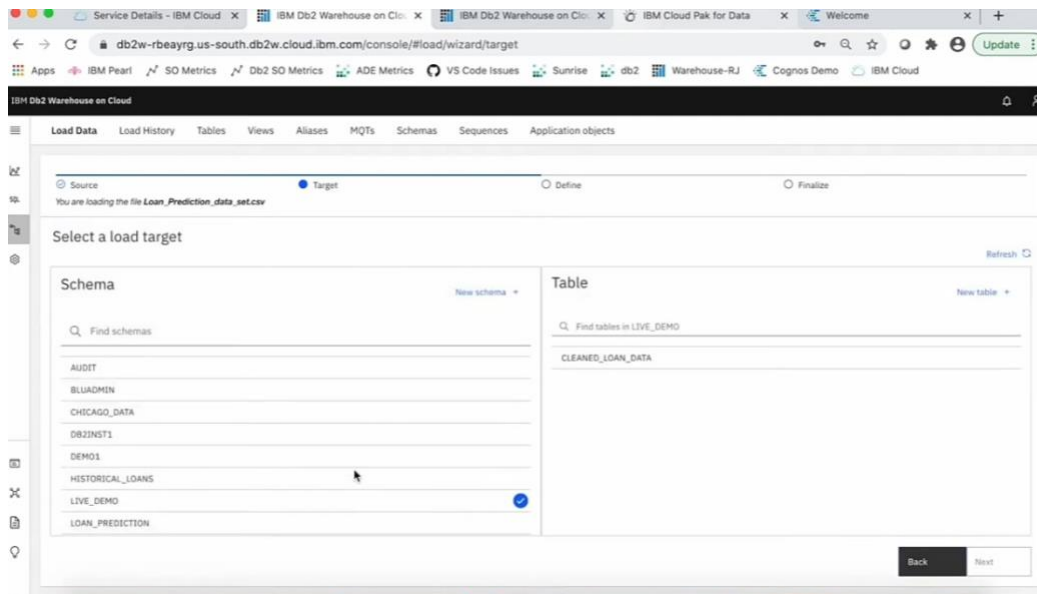


## 2. Data Transformation:

Clean and preprocess the extracted data. This includes tasks like data cleansing, data enrichment, and data normalization.

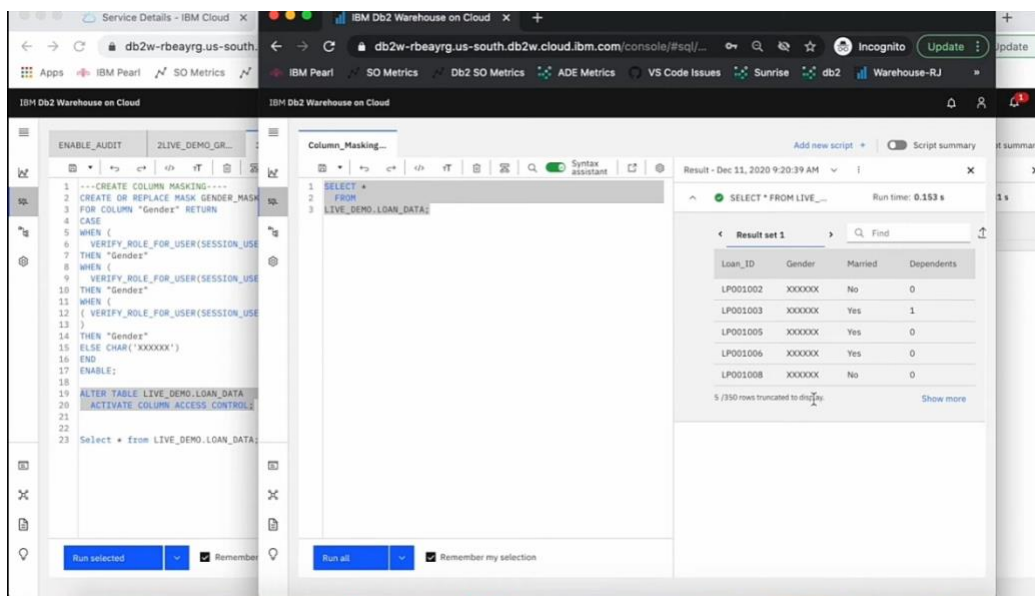
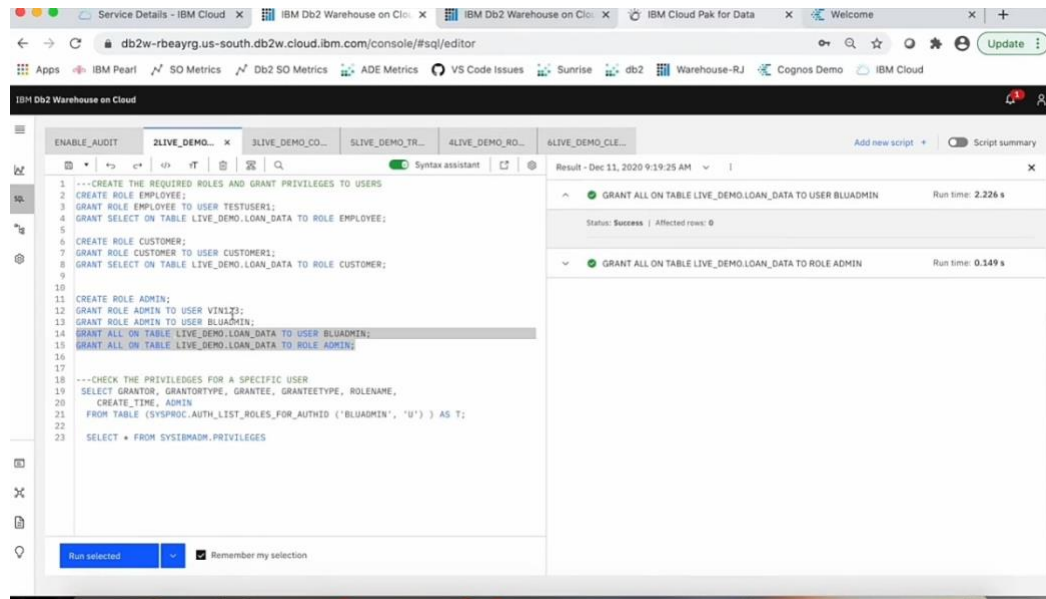
## 3.Data Loading:

Load the transformed data into your IBM Db2 Warehouse. Depending on your needs, you might use bulk loading, real-time streaming, or other loading methods.



#### 4. Data Exploration:

Provide data architects with tools and access to the data in Db2 Warehouse. SQL is a powerful tool for querying and analyzing data.



## **5. SQL Query Optimization:**

Ensure that your SQL queries are optimized for performance. This might involve creating indexes, tuning queries, and monitoring query performance.

## **6. Data Analysis Techniques:**

Support data architects in applying various data analysis techniques such as data mining, machine learning, and statistical analysis to derive insights from the data.

## **7. Data Governance:**

Implement data governance practices to ensure data quality, security, and compliance with relevant regulations.

## **8. Documentation:**

Document the ETL processes, data models, and analysis techniques for future reference and knowledge sharing.

## **How ETL works**

The easiest way to understand how ETL works is to understand what happens in each step of the process.

### **Extract**

During data extraction, raw data is copied or exported from source locations to a staging area. Data management teams can extract data from a variety of data sources, which can be structured or unstructured. Those sources include but are not limited to:

- SQL or [NoSQL](#) servers
- CRM and ERP systems
- Flat files
- Email
- Web pages

### **Transform**

In the staging area, the raw data undergoes data processing. Here, the data is transformed and consolidated for its intended analytical use case. This phase can involve the following tasks:

- Filtering, cleansing, de-duplicating, validating, and authenticating the data.

- Performing calculations, translations, or summarizations based on the raw data. This can include changing row and column headers for consistency, converting currencies or other units of measurement, editing text strings, and more.
- Conducting audits to ensure data quality and compliance
- Removing, encrypting, or protecting data governed by industry or governmental regulators
- Formatting the data into tables or joined tables to match the schema of the target data warehouse.

## Load

In this last step, the transformed data is moved from the staging area into a target data warehouse. Typically, this involves an initial loading of all data, followed by periodic loading of incremental data changes and, less often, full refreshes to erase and replace data in the warehouse.

For most organizations that use ETL, the process is automated, well-defined, continuous and batch-driven. Typically, ETL takes place during off-hours when traffic on the source systems and the data warehouse is at its lowest.

## ETL and other data integration methods

ETL and ELT are just two data integration methods, and there are other approaches that are also used to facilitate data integration workflows. Some of these include:

- **Change Data Capture (CDC)** identifies and captures only the source data that has changed and moves that data to the target system. CDC can be used to reduce the resources required during the ETL “extract” step; it can also be used independently to move data that has been transformed into a data lake or other repository in real time.
- **Data replication** copies changes in data sources in real time or in batches to a central database. [Data replication](#) is often listed as a data integration method. In fact, it is most often used to create backups for [disaster recovery](#).
- **Data virtualization** uses a software abstraction layer to create a unified, integrated, fully usable *view* of data—without physically copying, transforming or loading the source data to a target system. Data virtualization functionality enables an organization to create virtual data warehouses, data lakes and data marts from the same source data for data storage without the expense and complexity of building and managing separate platforms for each. While data virtualization can be used alongside ETL, it is increasingly seen as an alternative to ETL and to other physical data integration methods.
- **Stream Data Integration (SDI)** is just what it sounds like—it continuously consumes data streams in real time, transforms them, and loads them to a target system for analysis. The key word here is *continuously*. Instead of integrating snapshots of data extracted from sources at a given time, SDI integrates data constantly as it becomes available. SDI enables a data store for powering analytics, machine learning and real-time applications for improving customer experience, fraud detection and more.

## Create data model from SQL

In order to demonstrate how you can use the strategy defined in this document to build data model we will be using below mentioned queries.

/\*

Query to find all events and occurrences in one year

\*/

SELECT eve.eventname,count(\*) as total\_occurences

FROM Event eve

inner join

Date\_cal dat

on

eve.dateid = dat.dateid

where dat.caldate between '2022-01-01' and '2022-12-31'

group by eve.eventname

;/\*

Query to find all events on 30th Oct 2022 in Bangalore city

\*/

SELECT eve.eventname,

eve.starttime,

ven.venueName,

ven.venueCity,

ven.venueCityPop,

ven.venueState,

dat.day,

dat.holiday,

dat.caldate

FROM Event eve

inner join

Venue ven

on

eve.venueid = ven.venueid

inner join

Date\_cal dat

on

eve.dateid = dat.dateid

where ven.venuecity = 'Bangalore'

and dat.caldate = '2022-10-30'

;/\*

Query to find the total number of tickets sold, price paid, and commission paid for each event and then determine profit for each event.

\*/

Select sales\_event.eventname , sales\_event.tot\_Qtysold, sales\_event.tot\_Pricepaid -  
sales\_event.tot\_commission as tot\_profit

from(

select

sum(sal.Qtysold) as tot\_Qtysold,

sum(sal.Pricepaid) as tot\_Pricepaid,

sum(sal.commission) as tot\_commission,

eve.eventname

FROM

Sales sal

inner join

Event eve

on

sal.eventid = eve.eventid

group by eve.eventname

)sales\_event

order by 2 desc; /\*



Query to find least popular events in history

```
*/SELECT cat.catname,eve.eventname,sum(sal.Qtysold) as tot_tickets_sold
FROM Category cat
inner join
Event eve on
cat.catid = eve.catid
inner join
Sales sal on sal.eventid = eve.eventid
GROUP BY cat.catname,eve.eventname
ORDER BY 3 asc;
```

## SQL into Data Model – The Solution

The solution uses python packages with custom code to parse SQL and then generate **DBML output**. We will use the DBML output generated into the **web app to visualise it** and save as data model.

### Step 1:

#### Install the **sqlparse** package

```
pip install sqlparse
```

### Step 2: Install the sql-metadata package

```
pip install sql-metadata
```

### Step 3: Run the below python command

```
import sqlparse

from sql_metadata import Parser

strsql = ""

/*

Query to find all events and occurrences in one year

*/
```

```

SELECT eve.eventname,count(*) as total_occurences
FROM Event eve
inner join
Date_cal dat
on
eve.dateid = dat.dateid
where dat.caldate between '2022-01-01' and '2022-12-31'
group by eve.eventname;

```

/\*Query to find all events on 30th Oct 2022 in Bangalore city\*/

```

SELECT
eve.eventname,eve.starttime,ven.venueid,ven.venuecity,ven.venuecitypop,ven.venuestate,dat.day,dat.
holiday,dat.caldate
FROM Event eve
inner join Venue ven on
eve.venueid = ven.venueid inner join Date_cal dat
on
eve.dateid = dat.dateid
where ven.venuecity = 'Bangalore'
and dat.caldate = '2022-10-30';

```

/\*

Query to find the total number of tickets sold, price paid, and commission paid for each event and then determine profit for each event.

\*/

```

Select sales_event.eventname , sales_event.tot_Qtysold, sales_event.tot_Pricepaid -
sales_event.tot_commission as tot_profit
from
(
select
sum(sal.Qtysold) as tot_Qtysold,
sum(sal.Pricepaid) as tot_Pricepaid,

```

```
sum(sal.commission) as tot_commission,
```

```
eve.eventname
```

```
FROM
```

```
Sales sal
```

```
inner join
```

```
Event eve
```

```
on
```

```
sal.eventid = eve.eventid
```

```
group by eve.eventname
```

```
)sales_event
```

```
order by 2 desc;
```

```
/*
```

```
Query to find least popular events in history
```

```
*/
```

```
SELECT
```

```
cat.catname,
```

```
eve.eventname,
```

```
sum(sal.Qtysold) as tot_tickets_sold
```

```
FROM Category cat
```

```
inner join
```

```
Event eve on
```

```
cat.catid = eve.catid
```

```
inner join
```

```
Sales sal
```

```
on
```

```
sal.eventid = eve.eventid
```

```
GROUP BY cat.catname,eve.eventname
```

```
ORDER BY 3 asc;
```

```
""""
```

```
statements = sqlparse.split(strsql)
```

```

all_columns = list()

all_tables = list()

all_joins = ""

for stmt in statements:

    parser = Parser(stmt)

    for key,value in parser.columns_dict.items():

        if value is not None:

            all_columns.append(value)

    for j in parser.tables:

        all_tables.append(j.split())

    i=0

    join_cond=""

    if 'join' in parser.columns_dict:

        for j in parser.columns_dict['join']:

            if i%2 == 0:

                join_cond = join_cond + "" + "Ref: "+j

            else:

                join_cond = join_cond + " - " + j + "|"

            i = i + 1

        all_joins = all_joins + join_cond

all_joins_list = all_joins.split("|")

all_joins_list = list(set(all_joins_list))

all_joins = "|".join(all_joins_list)

all_columns = sum(all_columns, [])

all_columns = list(set(all_columns))

all_tables = sum(all_tables, [])

all_tables = list(set(all_tables))

i=0

table_def=""

```

```

for j in all_tables:
    table_struct=""
    table_struct = "Table "+j+" {\nrowid bigint"
    for i in all_columns:
        if i.startswith(j+'.') :
            table_struct = table_struct + "\n"+ i.split('.')[1] + " string "
    table_def = table_def + table_struct + "\n}\n"
for i in table_def[:-1].split("\n"):
    print (i)
for i in all_joins[1:].split("\n"):
    print (i)
-Change the SQL as per your requirement in the Python code.

```

## Step 4: Copy the DBML output generated

```

Table Date_cal {
rowid bigint
holiday string
day string
dateid string
caldate string
}
Table Event {
rowid bigint
starttime string
eventname string
venueid string
eventid string
dateid string
catid string
}

```

Table Category {

rowid bigint

catname string

catid string

}

Table Venue {

rowid bigint

venuecitypop string

venueid string

venuecity string

venuestate string

venueid string

}

Table Sales {

rowid bigint

commission string

Pricepaid string

eventid string

Qtysold string

}

Ref: Category.catid - Event.catid

Ref: Sales.eventid - Event.eventid

Ref: Event.dateid - Date\_cal.dateid

Ref: Event.venueid - Venue.venueid

## Step 5: Paste the DBML output

