

**IBM NAAN MUDHALVAN**  
**CLOUD APPLICATION DEVELOPMENT**  
**PROJECT**

***Data Warehousing with IBM Cloud Db2Warehouse***

**PHASE 5 : Documentation**

**Project Objective:**

The objective of this data warehousing project using IBM Db2 is to create a robust and scalable data warehouse solution to enable efficient data storage, retrieval, and analysis for the organization. This project aims to:

1. Centralize Data: Consolidate data from various sources, providing a single source of truth for the organization.
2. Improve Reporting and Analytics: Enhance the reporting and analytical capabilities by providing a structured and easily accessible data repository.
3. Enable Data-Driven Decision-Making: Empower decision-makers with timely, accurate, and actionable insights.
4. Optimize Performance: Ensure that data retrieval and processing are efficient and support real-time or near-real-time analytics.

**Design Thinking Process:**

The design thinking process for this project involves the following stages:

**1. Empathize:**

Understand the needs and pain points of stakeholders, including data analysts, business users, and IT teams. Conduct interviews and workshops to gather requirements.

**2. Define:**

Define the scope, goals, and key performance indicators (KPIs) for the data warehouse. Create user personas and use cases to guide design.

**3. Ideate:**

Brainstorm and generate creative ideas for the data warehouse architecture, data modeling, and integration strategies. Explore various possibilities for optimizing data storage and retrieval.

**4. Prototype:**

Create a prototype or proof of concept to visualize the data warehouse's structure and user interface. Test and iterate on the design based on user feedback.

**5. Test:**

Evaluate the prototype for usability, performance, and data quality. Identify and address any issues or bottlenecks.

**6. Implement:**

Develop and deploy the data warehouse using IBM Db2, following the best practices and design principles established during the previous stages.

**7. Iterate:**

Continuously improve the data warehouse based on user feedback and changing business needs. Implement enhancements and optimizations as required.

- Throughout the design thinking process, the project team should maintain close collaboration with stakeholders and ensure that the data warehouse solution aligns with the organization's objectives.
- This outline provides a framework for defining the project's objectives and the design thinking process that guides its development.


**DEVELOPMENT STEPS:**

In this part you will begin building your project.  
Start building the data warehouse using IBM Cloud Db2 Warehouse.

Define the schema and structure of the data warehouse tables. Identify data sources (e.g., CSV files, databases) and design a strategy to integrate them into the data warehouse.

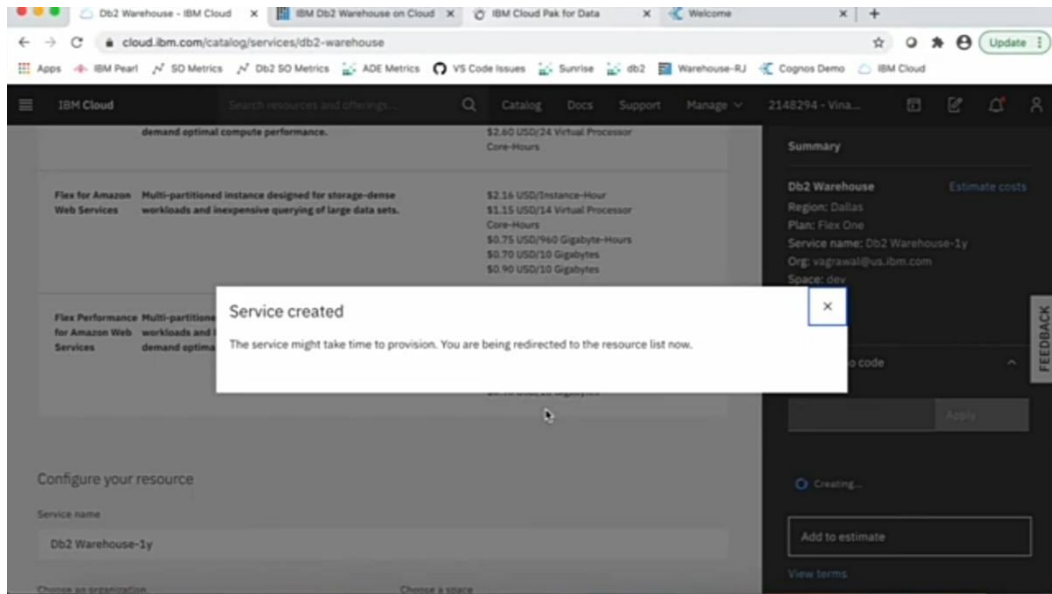
## 1. IBM Cloud Setup:

- Sign in to your IBM Cloud account or create one if you haven't already.

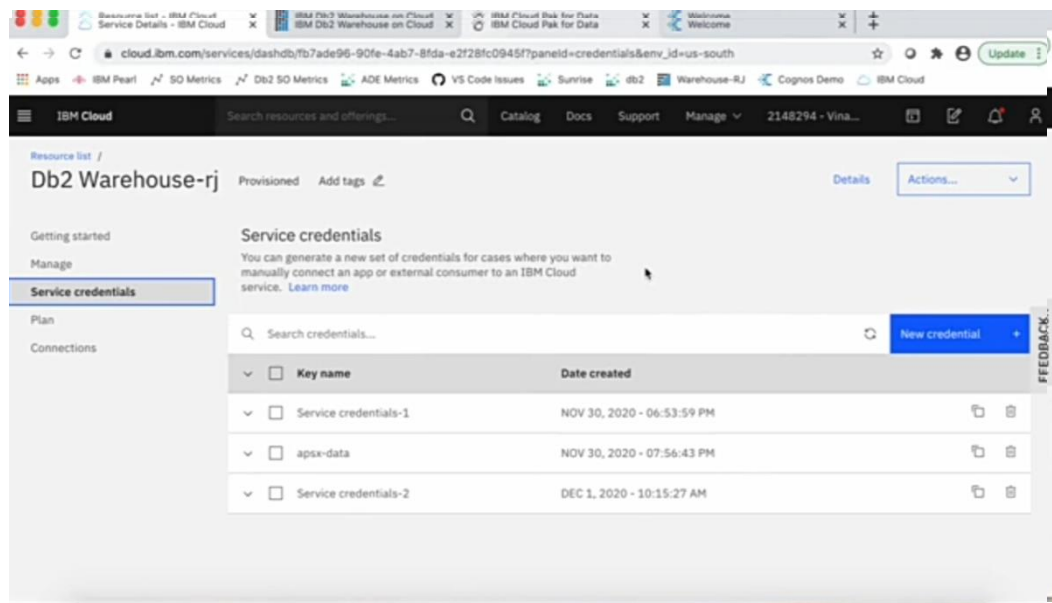


## 2. Provision Db2 Warehouse:

- In the IBM Cloud console, provision an instance of Db2 Warehouse.

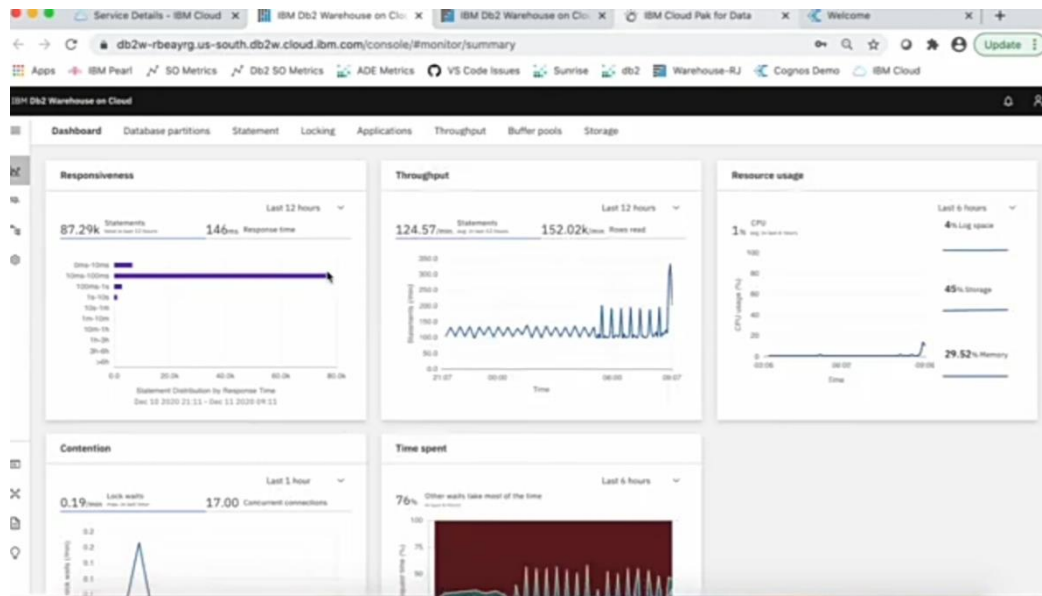
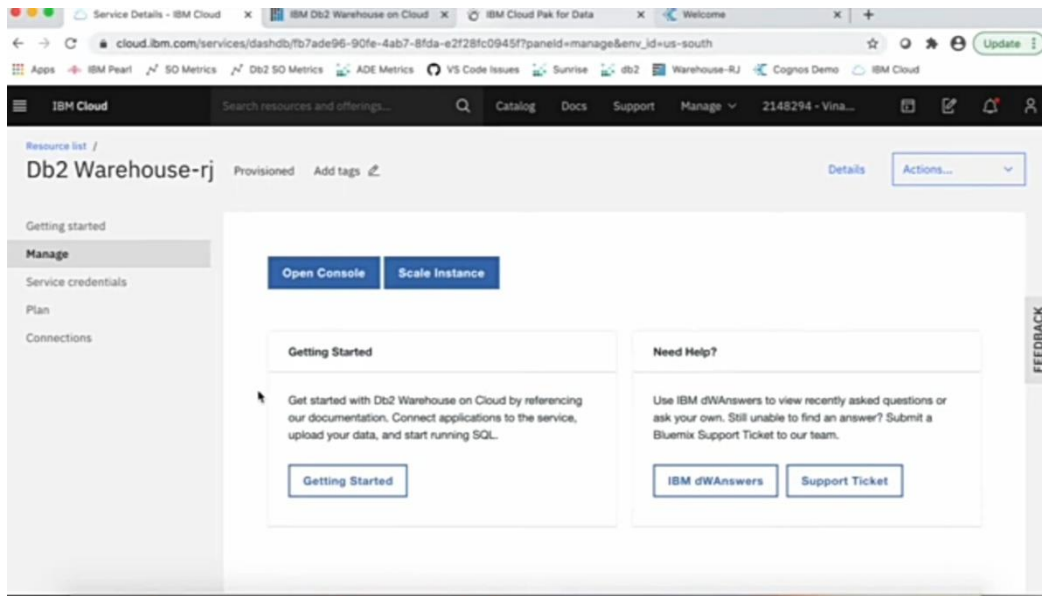


-can be viewed in resource list



### 3. Connect to Db2 Warehouse:

- Use a SQL client like Db2 Data Studio or command-line tools to connect to your Db2 Warehouse instance.



#### 4. Schema Definition:

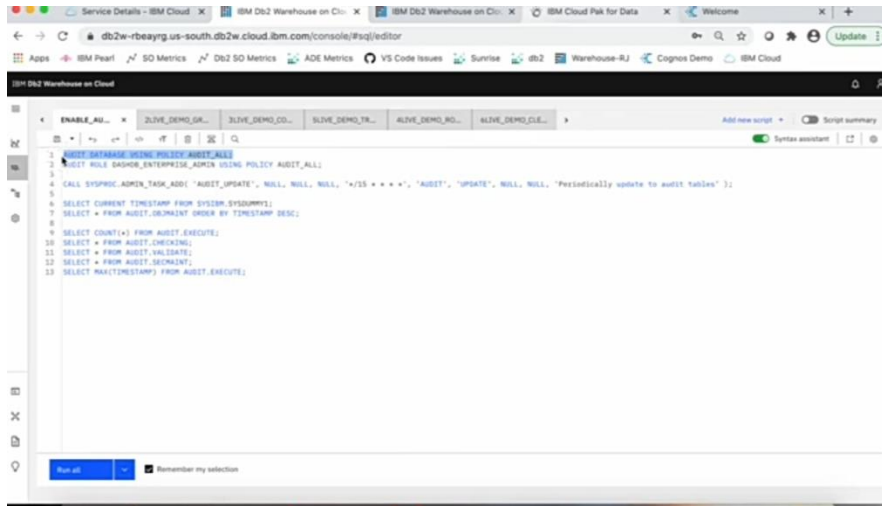
- Decide on the schema(s) for your data warehouse to logically organize data.
- Create a schema using SQL. For example:

```
```sql  
  
CREATE SCHEMA YOUR_SCHEMA;  
  
```
```

#### 5. Table Structure:

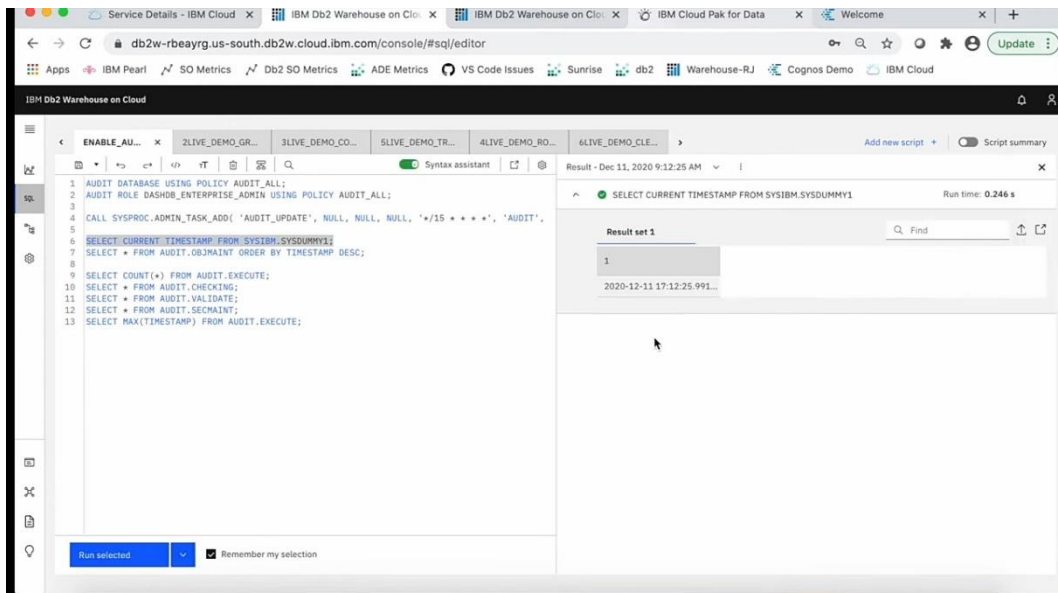
- Design the structure of your data warehouse tables. Determine which data attributes will be stored in each table.
- Create tables using SQL, specifying column names and data types. For example:

```
```sql  
  
CREATE TABLE YOUR_SCHEMA.YOUR_TABLE (  
    COLUMN1 DATATYPE,  
    COLUMN2 DATATYPE,  
    .  
....
```

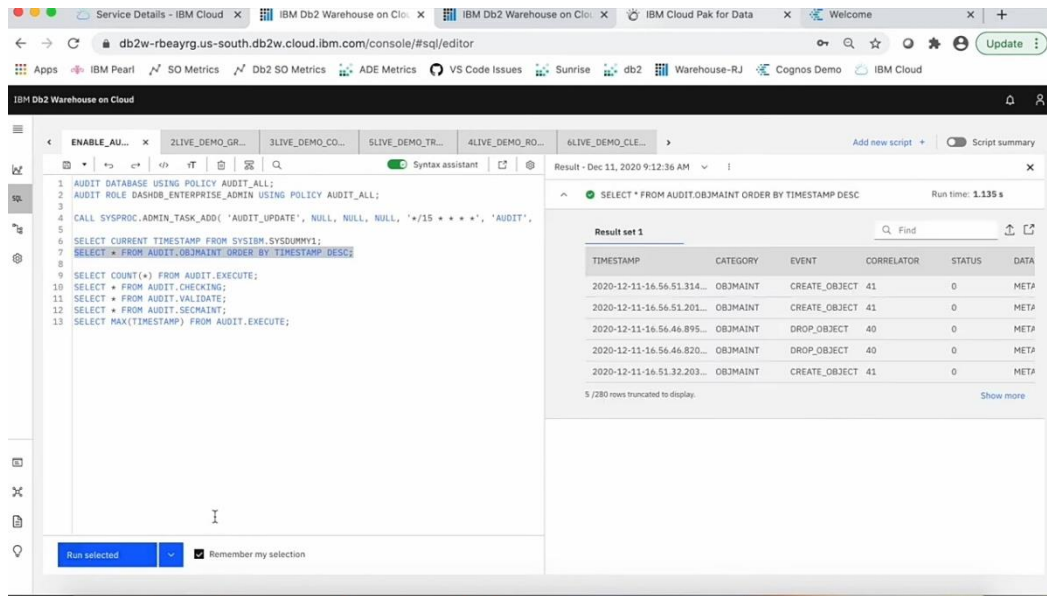


...

);



...



## 6. Data Sources Identification:

- Identify your data sources, which could include CSV files, other databases, APIs, etc.

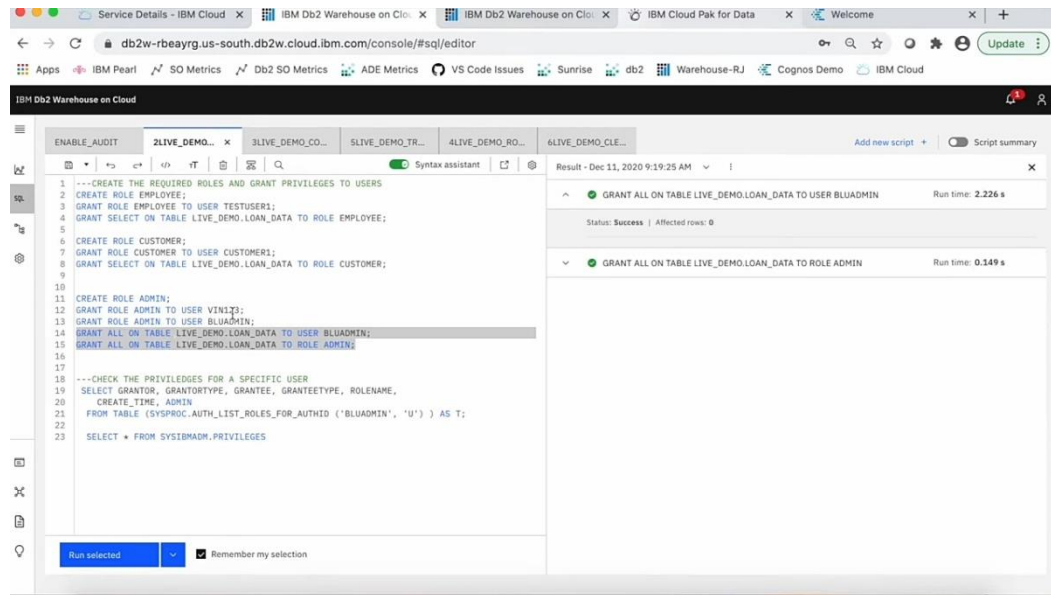
The screenshot shows the IBM Db2 Warehouse console interface, specifically the 'Load Data' section. The table 'LIVE\_DEMO.LOAN\_DATA' is displayed, showing a list of loan records. The table has columns: LOAN\_ID, GENDER, MARRIED, DEPENDENTS, EDUCATION, SELF EMPLOYED, APPLICANT, COAPPLICANT, LOAN AMOUNT, LOAN AMOUNT, CREDIT HISTORY, and PRIOR. The data is presented in a table format with 16 rows of loan records.

LOAN_ID	GENDER	MARRIED	DEPENDENTS	EDUCATION	SELF EMPLOYED	APPLICANT	COAPPLICANT	LOAN AMOUNT	LOAN AMOUNT	CREDIT HISTORY	PRIOR
1	Male	No	0	Graduate	No	58490.0	0.0	12800	36	1	Urt
2	Male	Yes	1	Graduate	No	45830.0	15080.0	12000	36	1	Ru
3	Male	Yes	0	Graduate	Yes	30000.0	0.0	6600	36	1	Urt
4	Male	Yes	0	Not Graduate	No	25830.0	23580.0	12000	36	1	Urt
5	Male	No	0	Graduate	No	60000.0	0.0	14100	36	1	Urt
6	Male	Yes	2	Graduate	Yes	54170.0	41960.0	26700	36	1	Urt
7	Male	Yes	0	Not Graduate	No	23330.0	15160.0	9500	36	1	Urt
8	Male	Yes	3	Graduate	No	30360.0	25040.0	15800	36	0	Sex
9	Male	Yes	2	Graduate	No	40060.0	15260.0	16800	36	1	Urt
10	Male	Yes	1	Graduate	No	128410.0	109680.0	34900	36	1	Sex
11	Male	Yes	2	Graduate	No	32000.0	7000.0	7000	36	1	Urt
12	Male	Yes	2	Graduate	No	25000.0	18400.0	10900	36	1	Urt
13	Male	Yes	2	Graduate	No	30730.0	81060.0	20000	36	1	Urt
14	Male	No	0	Graduate	No	18530.3	28400.0	11400	36	1	Ru
15	Male	Yes	2	Graduate	No	12990.0	10860.0	1700	12	1	Urt
16	Male	No	0	Graduate	No	49500.0	0.0	12500	36	1	Urt

## 7. Data Integration Strategy:

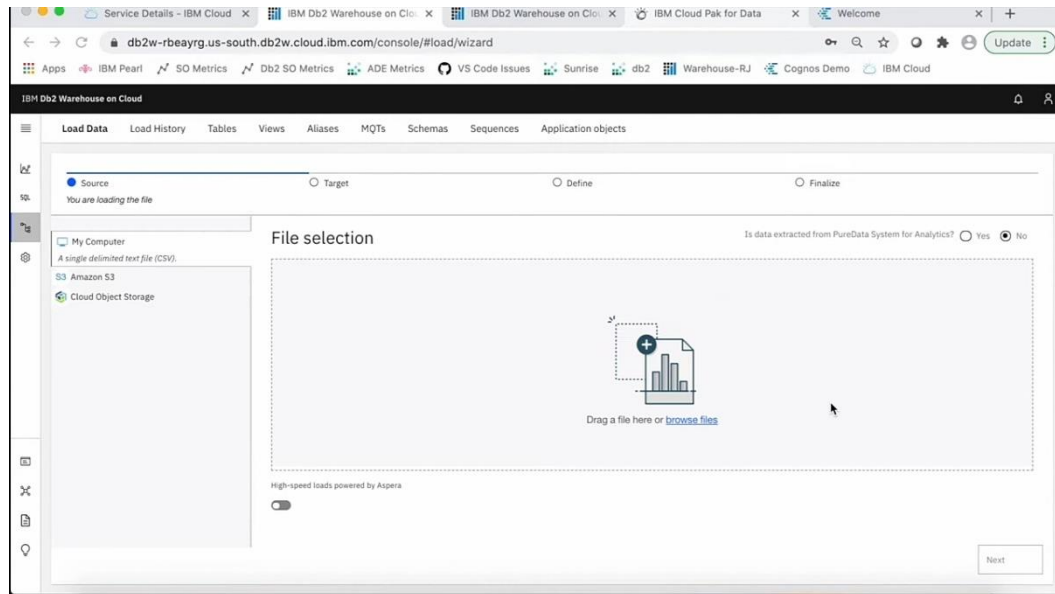


- Design a strategy to integrate data sources into your data warehouse. Consider using ETL (Extract, Transform, Load) tools or custom scripts.
- For CSV files, you might use Db2's `IMPORT` command to load data into your tables.
- For external databases, establish connections and design data transfer processes.

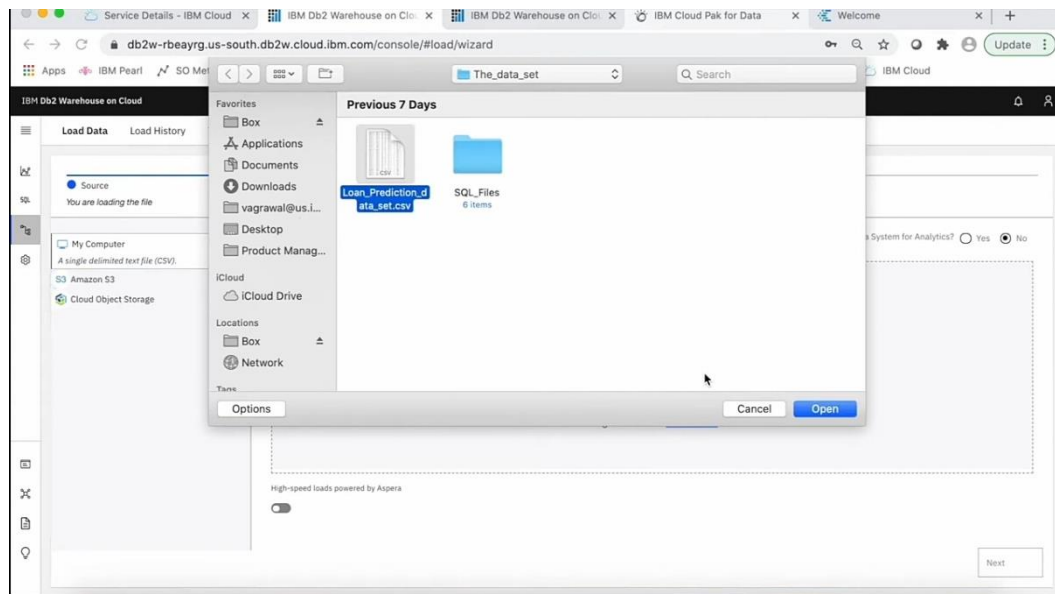


## 8. Data Loading:

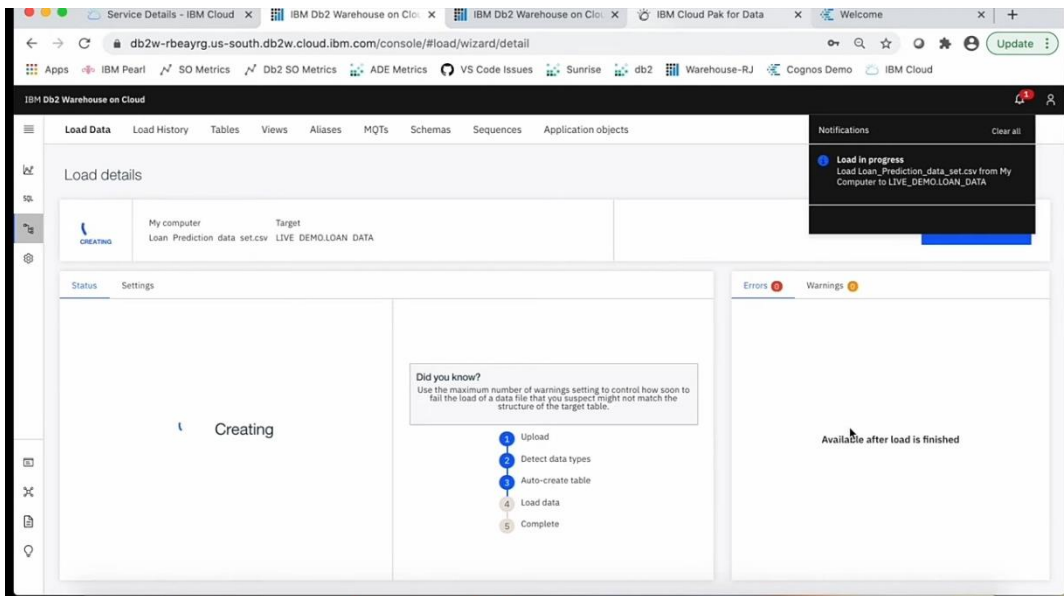
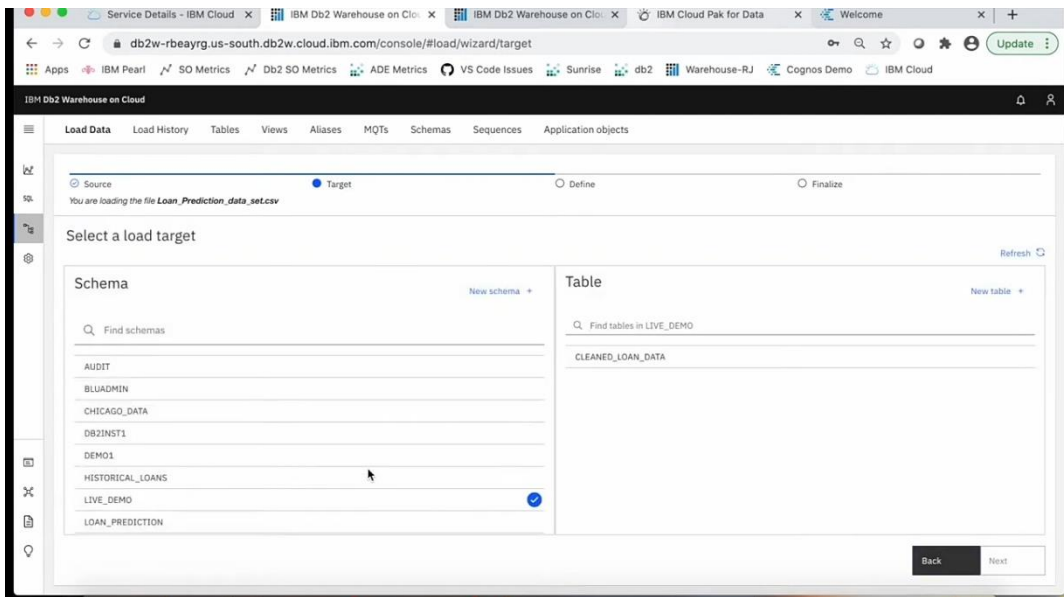
- Implement your data integration strategy to load data into your tables. This may involve running SQL scripts, using ETL tools, or other methods.



-add file



-load target file



## 9. Data Validation:

- Verify data integrity by running queries to validate that the data loaded correctly and is consistent with your expectations.

The screenshot shows the IBM Db2 Warehouse on Cloud console. The left pane displays a SQL script editor with the following content:

```
1 --CREATE THE REQUIRED ROLES AND GRANT PRIVILEGES TO US
2 CREATE ROLE EMPLOYEE;
3 GRANT ROLE EMPLOYEE TO USER TESTUSER1;
4 GRANT SELECT ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE EMPLOYEE;
5
6 CREATE ROLE CUSTOMER;
7 GRANT ROLE CUSTOMER TO USER CUSTOMER1;
8 GRANT SELECT ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE CUSTOMER;
9
10
11 CREATE ROLE ADMIN;
12 GRANT ROLE ADMIN TO USER VIN123;
13 GRANT ROLE ADMIN TO USER BLUADMIN;
14 GRANT ALL ON TABLE LIVE_DEMO.LOAN_DATA TO USER BLUADMIN;
15 GRANT ALL ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE ADMIN;
16
17
18 ---CHECK THE PRIVILEGES FOR A SPECIFIC USER
19 SELECT GRANTOR, GRANTORTYPE, GRANTEE, GRANTEETYPE, ROLENAME,
20        CREATE_TIME, ADMIN
21 FROM TABLE (SYSPROC.AUTH_LIST_ROLES_FOR_AUTHID ('BLUADMIN')) AS T;
22
23 SELECT * FROM SYSIBMADM.PRIVILEGES
```

The right pane shows the query results for the query `SELECT * FROM LIVE_DEMO.LOAN_DATA`. The results are displayed in a table with the following columns: Loan\_ID, Gender, Married, and Dependence. The table contains 8 rows of data, with the first 5 rows visible:

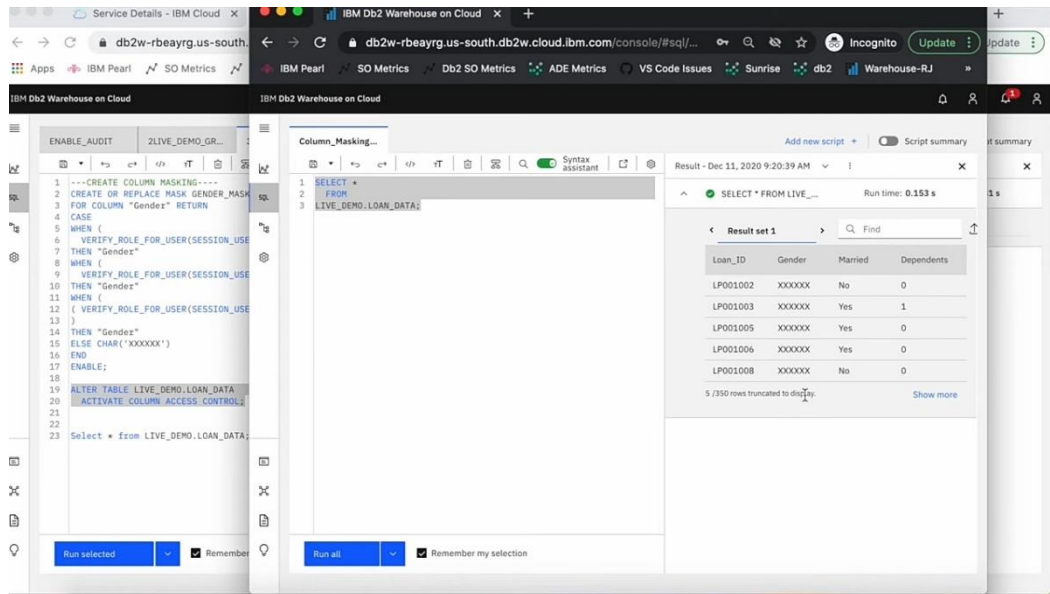
Loan_ID	Gender	Married	Dependence
LP001002	Male	No	0
LP001003	Male	Yes	1
LP001005	Male	Yes	0
LP001006	Male	Yes	0
LP001008	Male	No	0

The screenshot shows the IBM Db2 Warehouse on Cloud console. The left pane displays a SQL script editor with the following content:

```
1 --CREATE THE REQUIRED ROLES AND GRANT PRIVILEGES TO USERS
2 CREATE ROLE EMPLOYEE;
3 GRANT ROLE EMPLOYEE TO USER TESTUSER1;
4 GRANT SELECT ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE EMPLOYEE;
5
6 CREATE ROLE CUSTOMER;
7 GRANT ROLE CUSTOMER TO USER CUSTOMER1;
8 GRANT SELECT ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE CUSTOMER;
9
10
11 CREATE ROLE ADMIN;
12 GRANT ROLE ADMIN TO USER VIN123;
13 GRANT ROLE ADMIN TO USER BLUADMIN;
14 GRANT ALL ON TABLE LIVE_DEMO.LOAN_DATA TO USER BLUADMIN;
15 GRANT ALL ON TABLE LIVE_DEMO.LOAN_DATA TO ROLE ADMIN;
16
17
18 ---CHECK THE PRIVILEGES FOR A SPECIFIC USER
19 SELECT GRANTOR, GRANTORTYPE, GRANTEE, GRANTEETYPE, ROLENAME,
20        CREATE_TIME, ADMIN
21 FROM TABLE (SYSPROC.AUTH_LIST_ROLES_FOR_AUTHID ('BLUADMIN')) AS T;
22
23 SELECT * FROM SYSIBMADM.PRIVILEGES
```

The right pane shows the query results for the query `GRANT ALL ON TABLE LIVE_DEMO.LOAN_DATA TO USER BLUADMIN`. The results are displayed in a table with the following columns: Status, Affected rows, and Run time. The table contains 2 rows of data, with the first 2 rows visible:

Status	Affected rows	Run time
Success	0	2.226 s
Success	0	0.149 s



Certainly, here's an overview of the data warehouse structure, data integration strategies, ETL processes,

data exploration techniques, and how IBM Db2 Warehouse enables data architects to deliver actionable insights:

Data Warehouse Structure (IBM Db2 Warehouse):

### 1. Database Type:

IBM Db2 Warehouse is a relational database management system (RDBMS) optimized for analytical workloads, providing support for structured data.

### 2. Schema Design:

The data warehouse typically uses a star schema or snowflake schema for efficient querying. It includes fact tables for key business metrics and dimension tables for descriptive attributes.

### 3. Data Storage and Partitioning:

Data is stored in a columnar format for faster analytical queries. Partitioning is used to distribute data across multiple storage devices, optimizing query performance.

#### **4. Indexing Methods:**

IBM Db2 Warehouse supports various indexing methods, including B-tree, bitmap, and multidimensional clustering indexes, to accelerate data retrieval.

### **Data Integration Strategies:**

#### **1. Data Sources:**

Data can be integrated from various sources, including relational databases, external APIs, flat files, and streaming data sources.

#### **2. Data Extraction:**

Data is extracted from source systems using ETL processes, SQL queries, or data connectors provided by IBM Db2 Warehouse.

#### **3. Transformation Procedures:**

ETL processes involve data transformation to clean, structure, and format data for consistency and compatibility within the data warehouse.

#### **4. Loading Mechanisms:**

Data can be loaded into the warehouse using batch processes or near-real-time streaming to ensure data freshness.

### **ETL Processes (IBM Db2 Warehouse):**

#### **1. IBM DataStage:**

IBM DataStage, part of the IBM Information Server, is a popular ETL tool used to extract, transform, and load data into Db2 Warehouse. It provides a visual interface for designing ETL workflows.

#### **2. SQL Scripts:**

Data architects can also use SQL scripts to perform ETL operations, leveraging Db2 Warehouse's SQL capabilities for data transformation and loading.

### **3. Scheduling and Automation:**

ETL jobs can be scheduled and automated using IBM DataStage or other scheduling tools, ensuring that data is regularly updated in the warehouse.

In this part you will continue building your project. Continue building the data warehouse by implementing ETL processes and enabling data exploration.

Implement ETL processes to extract, transform, and load data into the data warehouse. Enable data architects to explore and analyze data within Db2 Warehouse using SQL queries and analysis techniques.

## **ETL**

ETL, which stands for extract, transform and load, is a data integration process that combines data from multiple data sources into a single, consistent data store that is loaded into a [data warehouse](#) or other target system.

As the databases grew in popularity in the 1970s, ETL was introduced as a process for integrating and loading data for computation and analysis, eventually becoming the primary method to process data for data warehousing projects.

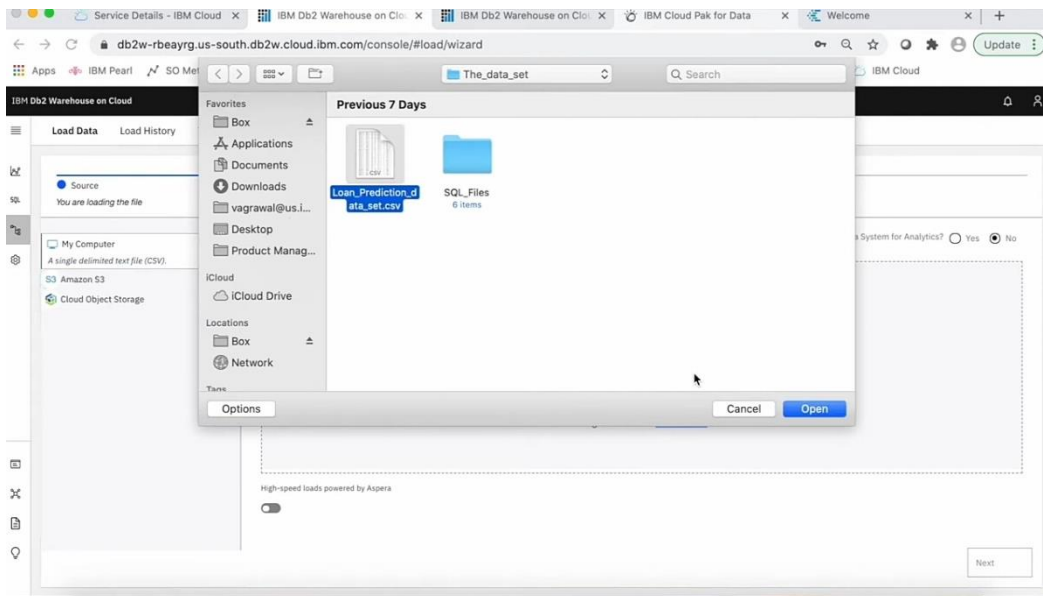
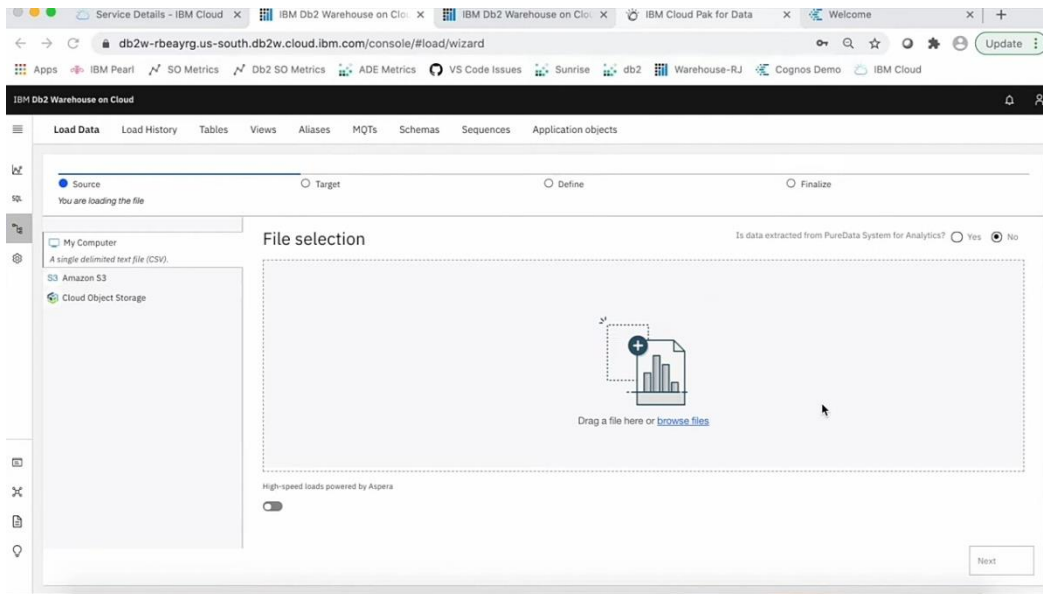
ETL provides the foundation for data analytics and machine learning workstreams. Through a series of business rules, ETL cleanses and organizes data in a way which addresses specific business intelligence needs, like monthly reporting, but it can also tackle more advanced analytics, which can improve back-end processes or end user experiences. ETL is often used by an organization to:

- Extract data from legacy systems
- Cleanse the data to improve data quality and establish consistency
- Load data into a target database

Here are some key steps you might consider:

### **1. Data Extraction:**

Extract data from various source systems. This could involve databases, files, APIs, or other data repositories.



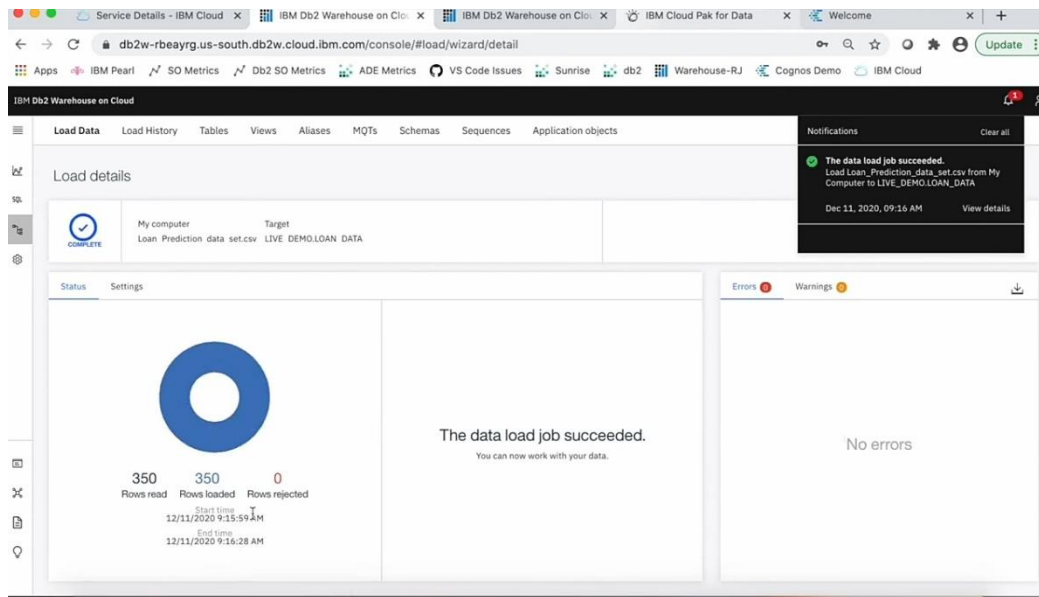
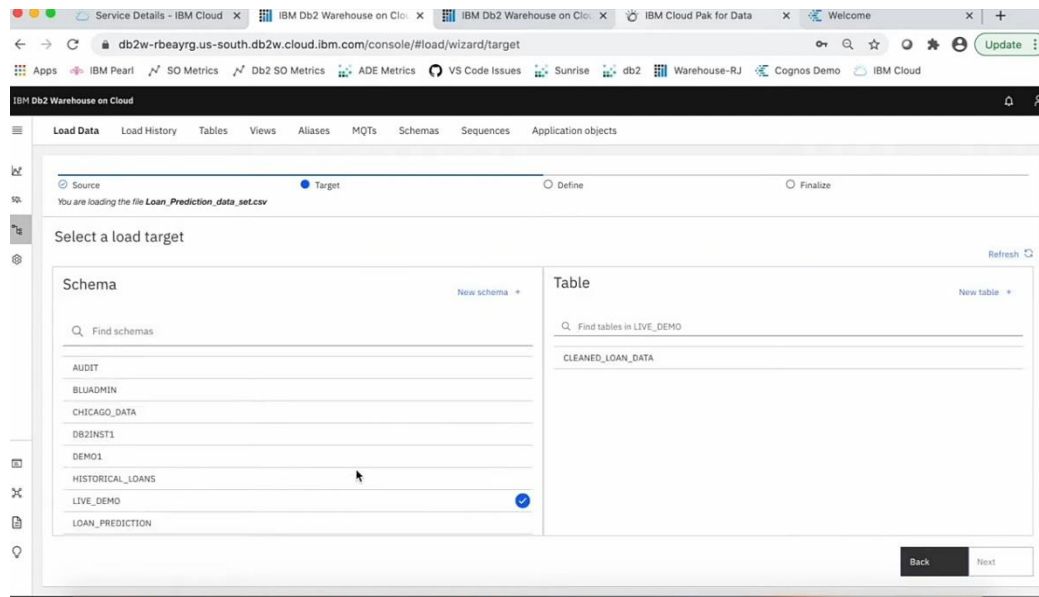
## 2. Data Transformation:

Clean and preprocess the extracted data. This includes tasks like data cleansing, data enrichment, and data normalization.

## 3.Data Loading:

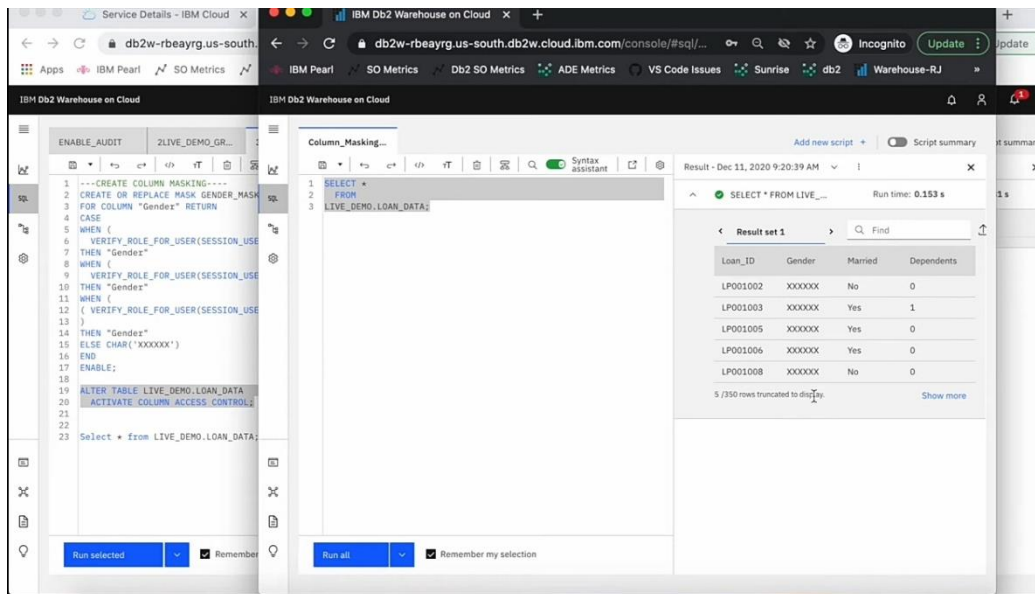
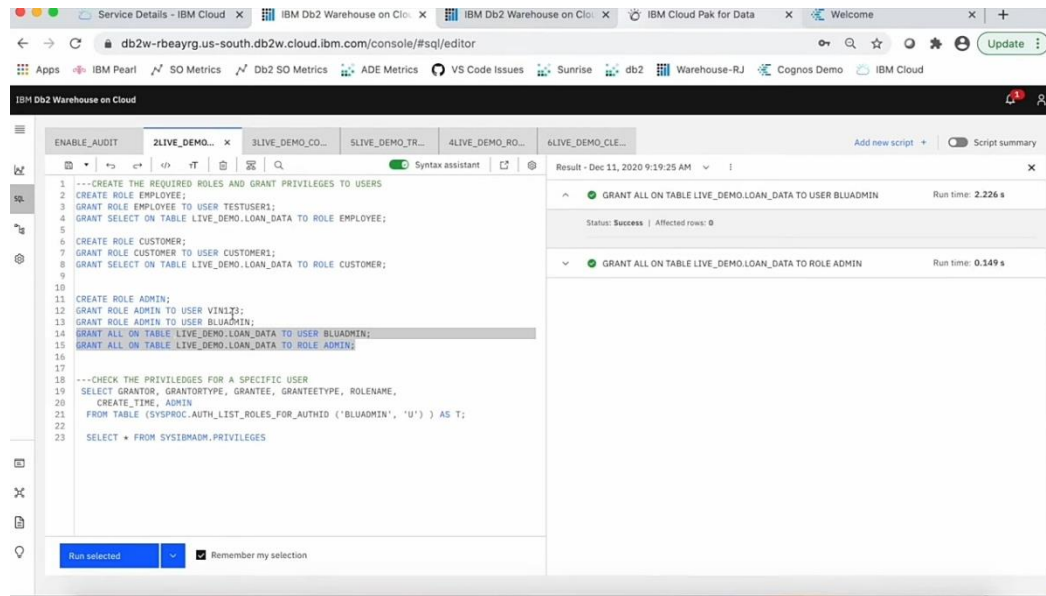


Load the transformed data into your IBM Db2 Warehouse. Depending on your needs, you might use bulk loading, real-time streaming, or other loading methods.



#### 4. Data Exploration:

Provide data architects with tools and access to the data in Db2 Warehouse. SQL is a powerful tool for querying and analyzing data.



## 5. SQL Query Optimization:

Ensure that your SQL queries are optimized for performance. This might involve creating indexes, tuning queries, and monitoring query performance.

## **6. Data Analysis Techniques:**

Support data architects in applying various data analysis techniques such as data mining, machine learning, and statistical analysis to derive insights from the data.

## **7. Data Governance:**

Implement data governance practices to ensure data quality, security, and compliance with relevant regulations.

## **8. Documentation:**

Document the ETL processes, data models, and analysis techniques for future reference and knowledge sharing.

## **How ETL works**

The easiest way to understand how ETL works is to understand what happens in each step of the process.

### **Extract**

During data extraction, raw data is copied or exported from source locations to a staging area. Data management teams can extract data from a variety of data sources, which can be structured or unstructured. Those sources include but are not limited to:

- SQL or [NoSQL](#) servers
- CRM and ERP systems
- Flat files
- Email
- Web pages

### **Transform**

In the staging area, the raw data undergoes data processing. Here, the data is transformed and consolidated for its intended analytical use case. This phase can involve the following tasks:

- Filtering, cleansing, de-duplicating, validating, and authenticating the data.

- Performing calculations, translations, or summarizations based on the raw data. This can include changing row and column headers for consistency, converting currencies or other units of measurement, editing text strings, and more.
- Conducting audits to ensure data quality and compliance
- Removing, encrypting, or protecting data governed by industry or governmental regulators
- Formatting the data into tables or joined tables to match the schema of the target data warehouse.

## Load

In this last step, the transformed data is moved from the staging area into a target data warehouse. Typically, this involves an initial loading of all data, followed by periodic loading of incremental data changes and, less often, full refreshes to erase and replace data in the warehouse.

For most organizations that use ETL, the process is automated, well-defined, continuous and batch-driven. Typically, ETL takes place during off-hours when traffic on the source systems and the data warehouse is at its lowest.

## ETL and other data integration methods

ETL and ELT are just two data integration methods, and there are other approaches that are also used to facilitate data integration workflows. Some of these include:

- **Change Data Capture (CDC)** identifies and captures only the source data that has changed and moves that data to the target system. CDC can be used to reduce the resources required during the ETL “extract” step; it can also be used independently to move data that has been transformed into a data lake or other repository in real time.
- **Data replication** copies changes in data sources in real time or in batches to a central database. [Data replication](#) is often listed as a data integration method. In fact, it is most often used to create backups for [disaster recovery](#).
- **Data virtualization** uses a software abstraction layer to create a unified, integrated, fully usable view of data—without physically copying, transforming or loading the source data to a target system. Data virtualization functionality enables an organization to create virtual data warehouses, data lakes and data marts from the same source data for data storage without the expense and complexity of building and managing separate platforms for each. While data virtualization can be used alongside ETL, it is increasingly seen as an alternative to ETL and to other physical data integration methods.
- **Stream Data Integration (SDI)** is just what it sounds like—it continuously consumes data streams in real time, transforms them, and loads them to a target system for analysis. The key word here is continuously. Instead of integrating snapshots of data extracted from sources at a given time, SDI integrates data constantly as it becomes available. SDI

enables a data store for powering analytics, machine learning and real-time applications for improving customer experience, fraud detection and more.

### Create data model from SQL

In order to demonstrate how you can use the strategy defined in this document to build data model we will be using below mentioned queries.

```
/*
```

Query to find all events and occurrences in one year

```
*/
```

```
SELECT eve.eventname,count(*) as total_occurences
```

```
FROM Event eve
```

```
inner join
```

```
Date_cal dat
```

```
on
```

```
eve.dateid = dat.dateid
```

```
where dat.caldate between '2022-01-01' and '2022-12-31'
```

```
group by eve.eventname
```

```
;/*
```

Query to find all events on 30th Oct 2022 in Bangalore city

```
*/
```

```
SELECT eve.eventname,
```

```
eve.starttime,
```

```
ven.venueName,
```

```
ven.venueCity,
```

```
ven.venueCityPop,
```

```
ven.venueState,
```

```

dat.day,
dat.holiday,
dat.caldate
FROM Event eve
inner join
Venue ven
on
eve.venueid = ven.venueid
inner join
Date_cal dat
on
eve.dateid = dat.dateid
where ven.venuecity = 'Bangalore'
and dat.caldate = '2022-10-30'
; /*

```

Query to find the total number of tickets sold, price paid, and commission paid for each event and then determine profit for each event.

```

*/
Select sales_event.eventname , sales_event.tot_Qtysold, sales_event.tot_Pricepaid -
sales_event.tot_commission as tot_profit
from(
select
sum(sal.Qtysold) as tot_Qtysold,
sum(sal.Pricepaid) as tot_Pricepaid,
sum(sal.commission) as tot_commission,
eve.eventname
FROM
Sales sal
inner join

```

```
Event eve
on
sal.eventid = eve.eventid
group by eve.eventname
)sales_event
order by 2 desc;/*
```

Query to find least popular events in history

```
*/SELECT cat.catname,eve.eventname,sum(sal.Qtysold) as tot_tickets_sold
FROM Category cat
inner join
Event eve on
cat.catid = eve.catid
inner join
Sales sal on sal.eventid = eve.eventid
GROUP BY cat.catname,eve.eventname
ORDER BY 3 asc;
```

### SQL into Data Model – The Solution

The solution uses python packages with custom code to parse SQL and then generate **DBML output**. We will use the DBML output generated into the **web app to visualise it** and save as data model.

#### Step 1:

##### Install the **sqlparse package**

```
pip install sqlparse
```

#### Step 2: Install the sql-metadata package

```
pip install sql-metadata
```

Step 3: Run the below python command

```
import sqlparse
```

```
from sql_metadata import Parser
```

```
strsql = ""
```

```
/*
```

```
Query to find all events and occurrences in one year
```

```
*/
```

```
SELECT eve.eventname,count(*) as total_occurences
```

```
FROM Event eve
```

```
inner join
```

```
Date_cal dat
```

```
on
```

```
eve.dateid = dat.dateid
```

```
where dat.caldate between '2022-01-01' and '2022-12-31'
```

```
group by eve.eventname;
```

```
/*Query to find all events on 30th Oct 2022 in Bangalore city*/
```

```
SELECT
```

```
eve.eventname,eve.starttime,ven.venueid,ven.venuecity,ven.venuecitypop,ven.venuestate,dat.  
day,dat.holiday,dat.caldate
```

```
FROM Event eve
```

```
inner join Venue ven on
```

```
eve.venueid = ven.venueid inner join Date_cal dat
```

```
on
```

```
eve.dateid = dat.dateid
```

```
where ven.venuecity = 'Bangalore'
```

```
and dat.caldate = '2022-10-30';
```

```
/*
```



Query to find the total number of tickets sold, price paid, and commission paid for each event and then determine profit for each event.

\*/

```
Select sales_event.eventname , sales_event.tot_Qtysold, sales_event.tot_Pricepaid -  
sales_event.tot_commission as tot_profit
```

```
from
```

```
(
```

```
select
```

```
sum(sal.Qtysold) as tot_Qtysold,
```

```
sum(sal.Pricepaid) as tot_Pricepaid,
```

```
sum(sal.commission) as tot_commission,
```

```
eve.eventname
```

```
FROM
```

```
Sales sal
```

```
inner join
```

```
Event eve
```

```
on
```

```
sal.eventid = eve.eventid
```

```
group by eve.eventname
```

```
)sales_event
```

```
order by 2 desc;
```

```
/*
```

Query to find least popular events in history

\*/

```
SELECT
```

```
cat.catname,
```

```
eve.eventname,
```

```
sum(sal.Qtysold) as tot_tickets_sold
```

```
FROM Category cat
```

```

inner join
Event eve on
cat.catid = eve.catid
inner join
Sales sal
on
sal.eventid = eve.eventid
GROUP BY cat.catname,eve.eventname
ORDER BY 3 asc;
"""

```

```

statements = sqlparse.split(strsql)
all_columns = list()
all_tables = list()
all_joins = ""

```

```

for stmt in statements:
    parser = Parser(stmt)
    for key,value in parser.columns_dict.items():
        if value is not None:
            all_columns.append(value)
    for j in parser.tables:
        all_tables.append(j.split())
    i=0
    join_cond=""
    if 'join' in parser.columns_dict:
        for j in parser.columns_dict['join']:
            if i%2 == 0:
                join_cond = join_cond + "" + "Ref: "+j

```

```

        else:
            join_cond = join_cond + " - " + j + "|"
            i = i + 1
        all_joins = all_joins + join_cond
    all_joins_list = all_joins.split("|")
    all_joins_list = list(set(all_joins_list))
    all_joins = '|'.join(all_joins_list)
    all_columns = sum(all_columns, [])
    all_columns = list(set(all_columns))
    all_tables = sum(all_tables, [])
    all_tables = list(set(all_tables))
    i=0
    table_def=""
    for j in all_tables:
        table_struct=""
        table_struct = "Table "+j+" {\nrowid bigint"
        for i in all_columns:
            if i.startswith(j+'.') :
                table_struct = table_struct + "\n" + i.split('.')[0] + " string "
        table_def = table_def + table_struct + "\n}\n|"
    for i in table_def[:-1].split("|"):
        print (i)
    for i in all_joins[1:].split("|"):
        print (i)

```

-Change the SQL as per your requirement in the Python code.

Step 4: Copy the DBML output generated

```

Table Date_cal {
rowid bigint

```

holiday string

day string

dateid string

caldate string

}

Table Event {

rowid bigint

starttime string

eventname string

venueid string

eventid string

dateid string

catid string

}

Table Category {

rowid bigint

catname string

catid string

}

Table Venue {

rowid bigint

venuecitypop string

venueid string

venuecity string

venuestate string

venue name string

}

```

Table Sales {
rowid bigint
commission string
Pricepaid string
eventid string
Qtysold string
}

```

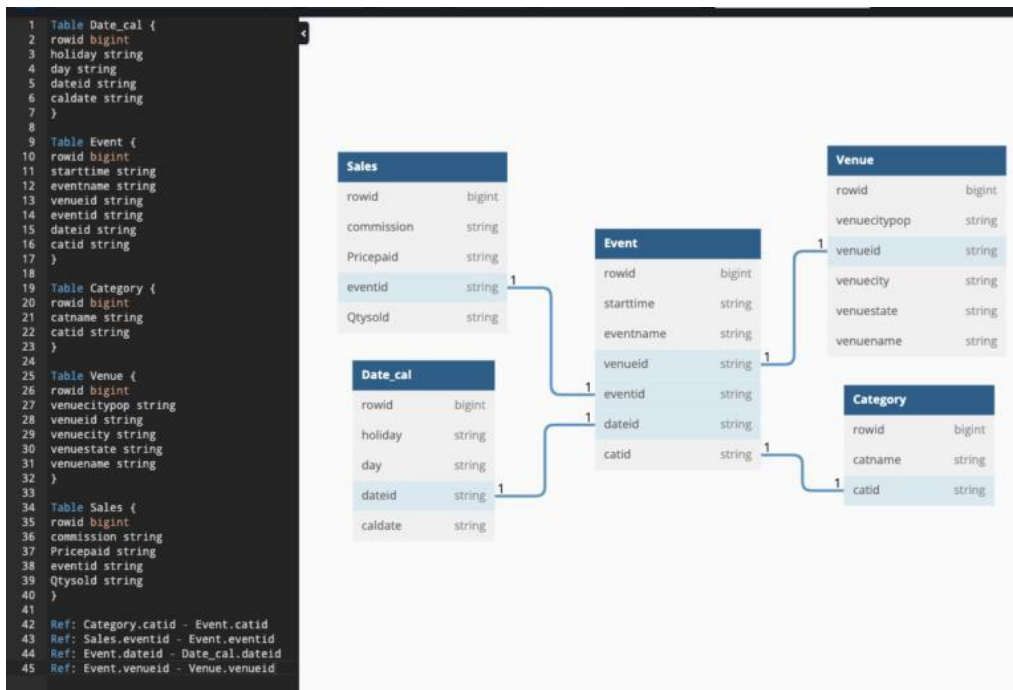
Ref: Category.catid - Event.catid

Ref: Sales.eventid - Event.eventid

Ref: Event.dateid - Date\_cal.dateid

Ref: Event.venueid - Venue.venueid

Step 5: Paste the DBML output



**Data Exploration Techniques (IBM Db2 Warehouse):**

## **1.SQL Queries:**

Data architects and analysts can use SQL queries to explore and analyze data stored in IBM Db2 Warehouse. SQL's powerful analytical functions enable complex data exploration.

## **2.IBM Cognos Analytics:**

IBM Cognos Analytics is a business intelligence and data visualization tool that integrates seamlessly with Db2 Warehouse. It allows for the creation of interactive dashboards and reports.

## **3.Jupyter Notebooks:**

Data scientists and analysts can use Jupyter notebooks with Python or R to perform in-depth data analysis and visualization using data from Db2 Warehouse.

## **Actionable Insights (IBM Db2 Warehouse):**

### **1. Real-time Analytics:**

Db2 Warehouse supports near-real-time data updates, enabling data architects to provide real-time or near-real-time analytics to stakeholders.

### **2. Advanced Analytics:**

With Db2 Warehouse, data architects can incorporate advanced analytics, machine learning models, and statistical analysis into the data exploration process, delivering deeper insights.

### **3. Scalability and Performance:**

Db2 Warehouse's scalable architecture ensures that it can handle large volumes of data and complex queries, facilitating the delivery of actionable insights without performance bottlenecks.

### **4. Integration with IBM Watson:**

For organizations using IBM Watson, integration with Db2 Warehouse allows for AI-powered insights and predictions, enhancing the value of the data warehouse.

In summary, IBM Db2 Warehouse offers a robust data warehousing solution with a structured data warehouse, efficient data integration, ETL capabilities, and powerful data exploration techniques. It empowers data architects to deliver actionable insights by providing real-time analytics, advanced analytics, and seamless integration with BI and AI tools.